

```

import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import os
import glob
import cv2
from PIL import Image
from torch.utils.data import Dataset, DataLoader
from albumentations import HorizontalFlip, VerticalFlip, Rotate
import tqdm
import torch.nn.functional as F
import matplotlib.image as mpimg

```

```

torch.manual_seed(42)
np.random.seed(42)
torch.cuda.manual_seed_all(42)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False

```

```
!nvidia-smi
```

```
Fri Jun  3 02:37:05 2022
```

```

+-----+
| NVIDIA-SMI 460.32.03      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+-----+-----+-----+-----+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                    |    MIG M.     |
+-----+-----+-----+-----+-----+-----+
|   0   Tesla T4              Off      | 00000000:00:04:0 Off |                    0 |
| N/A   47C    P8             9W / 70W |  0MiB / 15109MiB |      0%      Default |
|                                       |                    | N/A           |
+-----+-----+-----+-----+-----+

```

```

+-----+
| Processes:                                     |
|  GPU   GI    CI          PID    Type    Process name                  GPU Memory |
|          ID    ID                                   Usage          |
+-----+-----+-----+-----+-----+
| No running processes found                  |
+-----+

```

```
print(torch.cuda.is_available())
```

```
True
```

```
size = (256, 256)
```

```
class LoadData(Dataset):
```

```
def __init__(self, images_path, masks_path):
    super().__init__()

    self.images_path = images_path
    self.masks_path = masks_path
    self.len = len(images_path)
    self.transform = transforms.Resize(size)

def __getitem__(self, idx):
    img = Image.open(self.images_path[idx])
    img = self.transform(img)
    img = np.transpose(img, (2, 0, 1))
    img = img/255.0
    img = torch.tensor(img)

    mask = Image.open(self.masks_path[idx]).convert('L')
    mask = self.transform(mask)
    mask = np.expand_dims(mask, axis=0)
    mask = mask/255.0
    mask = torch.tensor(mask)

    return img, mask

def __len__(self):
    return self.len
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
X = sorted(glob.glob('/content/drive/MyDrive/satellite/images/*.jpg'))
y = sorted(glob.glob('/content/drive/MyDrive/satellite/masks/*.jpg'))
```

```
len(X)
```

2841

```
len(y)
```

2841

```
train_X = X[:2800]
train_y = y[:2800]
```

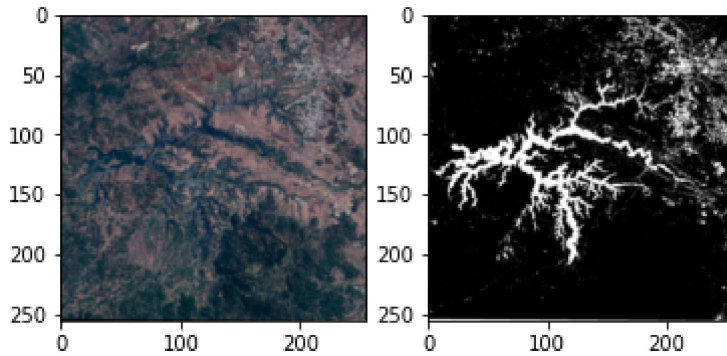
```
valid_X = X[2800:]
valid_y = y[2800:]
```

```
train_dataset = LoadData(train_X, train_y)
valid_dataset = LoadData(valid_X, valid_y)
```

```
img, mask = train_dataset[0]
```

```
f, axarr = plt.subplots(1,2)
axarr[1].imshow(np.squeeze(mask.numpy()), cmap='gray')
axarr[0].imshow(np.transpose(img.numpy(), (1,2,0)))
```

<matplotlib.image.AxesImage at 0x7f306ae8ea90>



```
img.shape
```

```
torch.Size([3, 256, 256])
```

```
class conv(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()

        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(out_channels)

        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3, padding=1)
        self.bn2 = nn.BatchNorm2d(out_channels)

        self.relu = nn.ReLU()

    def forward(self, images):
        x = self.conv1(images)
        x = self.bn1(x)
        x = self.relu(x)

        x = self.conv2(x)
        x = self.bn2(x)
        x = self.relu(x)

        return x
```

```
class encoder(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()
```

```

        self.conv = conv(in_channels, out_channels)
        self.pool = nn.MaxPool2d((2,2))

    def forward(self, images):
        x = self.conv(images)
        p = self.pool(x)

        return x, p

class decoder(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()

        self.upconv = nn.ConvTranspose2d(in_channels, out_channels, kernel_size=2, stride=
        self.conv = conv(out_channels * 2, out_channels)

    def forward(self, images, prev):
        x = self.upconv(images)
        x = torch.cat([x, prev], axis=1)
        x = self.conv(x)

        return x

class UNet(nn.Module):
    def __init__(self):
        super().__init__()

        self.e1 = encoder(3, 64)
        self.e2 = encoder(64, 128)
        self.e3 = encoder(128, 256)
        self.e4 = encoder(256, 512)

        self.b = conv(512, 1024)

        self.d1 = decoder(1024, 512)
        self.d2 = decoder(512, 256)
        self.d3 = decoder(256, 128)
        self.d4 = decoder(128, 64)

        self.output = nn.Conv2d(64, 1, kernel_size=1, padding=0)

    def forward(self, images):
        x1, p1 = self.e1(images)
        x2, p2 = self.e2(p1)
        x3, p3 = self.e3(p2)
        x4, p4 = self.e4(p3)

        b = self.b(p4)

        d1 = self.d1(b, x4)
        d2 = self.d2(d1, x3)
        d3 = self.d3(d2, x2)
        d4 = self.d4(d3, x1)

```

```

        output_mask = self.output(d4)
        output_mask = torch.sigmoid(output_mask)

    return output_mask

batch_size = 8
num_epochs = 12
lr = 1e-4
checkpoint_path = "./checkpoint.pth"

train_loader = DataLoader(
    dataset=train_dataset,
    batch_size=batch_size,
    shuffle=True,
    num_workers=2,
)

valid_loader = DataLoader(
    dataset=valid_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=2,
)

device = torch.device('cuda')
model = UNet()
model = model.to(device)

class DiceBCELoss(nn.Module):
    def __init__(self, weight=None, size_average=True):
        super(DiceBCELoss, self).__init__()

    def forward(self, inputs, targets, smooth=1):
        inputs = inputs.view(-1)
        targets = targets.view(-1)

        intersection = (inputs * targets).sum()
        dice_score = (2.*intersection + smooth)/(inputs.sum() + targets.sum() + smooth)
        dice_loss = 1 - dice_score
        loss = torch.nn.BCELoss()
        BCE = loss(inputs, targets)
        Dice_BCE = BCE + dice_loss

        return Dice_BCE

optimizer = torch.optim.Adam(model.parameters(), lr=lr)
loss_fn = DiceBCELoss()

def train_model(model, loader, optimizer, loss_fn, device):

```

```

epoch_loss = 0.0

model.train()
for x, y in loader:
    x = x.to(device, dtype=torch.float32)
    y = y.to(device, dtype=torch.float32)

    optimizer.zero_grad()
    y_pred = model(x)
    loss = loss_fn(y_pred, y)
    loss.backward()
    optimizer.step()
    epoch_loss += loss.item()

epoch_loss = epoch_loss/len(loader)
return epoch_loss

def evaluate(model, loader, loss_fn, device):
    epoch_loss = 0.0

    model.eval()
    with torch.no_grad():
        for x, y in loader:
            x = x.to(device, dtype=torch.float32)
            y = y.to(device, dtype=torch.float32)
#             x,y=x.type(torch.FloatTensor),y.type(torch.FloatTensor)

            y_pred = model(x)
            loss = loss_fn(y_pred, y)
            epoch_loss += loss.item()

    epoch_loss = epoch_loss/len(loader)
    return epoch_loss

train = []
valid = []

best_valid_loss = float("inf")

for epoch in range(num_epochs):
    train_loss = train_model(model, train_loader, optimizer, loss_fn, device)
    valid_loss = evaluate(model, valid_loader, loss_fn, device)

    train.append(train_loss)
    valid.append(valid_loss)

    if valid_loss < best_valid_loss:
        data_str = f"Valid loss improved from {best_valid_loss:2.4f} to {valid_loss:2.4f}"
        print(data_str)

        best_valid_loss = valid_loss
        torch.save(model.state_dict(), checkpoint_path)

```

```
data_str = f'Epoch: {epoch+1:02}\n'  
data_str += f'\tTrain Loss: {train_loss:.3f}\n'  
data_str += f'\t Val. Loss: {valid_loss:.3f}\n'  
print(data_str)
```

Valid loss improved from inf to 0.6455. Saving checkpoint: ./checkpoint.pth  
Epoch: 01

Train Loss: 0.843  
Val. Loss: 0.645

Valid loss improved from 0.6455 to 0.5781. Saving checkpoint: ./checkpoint.pth  
Epoch: 02

Train Loss: 0.732  
Val. Loss: 0.578

Epoch: 03

Train Loss: 0.659  
Val. Loss: 0.595

Valid loss improved from 0.5781 to 0.5382. Saving checkpoint: ./checkpoint.pth  
Epoch: 04

Train Loss: 0.605  
Val. Loss: 0.538

Valid loss improved from 0.5382 to 0.5331. Saving checkpoint: ./checkpoint.pth  
Epoch: 05

Train Loss: 0.573  
Val. Loss: 0.533

Valid loss improved from 0.5331 to 0.5278. Saving checkpoint: ./checkpoint.pth  
Epoch: 06

Train Loss: 0.556  
Val. Loss: 0.528

Valid loss improved from 0.5278 to 0.5050. Saving checkpoint: ./checkpoint.pth  
Epoch: 07

Train Loss: 0.533  
Val. Loss: 0.505

Valid loss improved from 0.5050 to 0.5017. Saving checkpoint: ./checkpoint.pth  
Epoch: 08

Train Loss: 0.520  
Val. Loss: 0.502

Valid loss improved from 0.5017 to 0.4902. Saving checkpoint: ./checkpoint.pth  
Epoch: 09

Train Loss: 0.513  
Val. Loss: 0.490

Epoch: 10

Train Loss: 0.491  
Val. Loss: 0.512

Valid loss improved from 0.4902 to 0.4874. Saving checkpoint: ./checkpoint.pth  
Epoch: 11

Train Loss: 0.482  
Val. Loss: 0.487

Valid loss improved from 0.4874 to 0.4770. Saving checkpoint: ./checkpoint.pth


Epoch: 12

Train Loss: 0.479

Val. Loss: 0.477

```
plt.plot(range(0,12), train, label='train loss')
plt.plot(range(0,12), valid, label='valid loss')
plt.legend()
```

&lt;matplotlib.legend.Legend at 0x7f30685e6750&gt;



```
m = UNet()
m.load_state_dict(torch.load(checkpoint_path))
m = m.to(device)
```

```
transform = transforms.ToPILImage()
```

```
pred = []
```

```
for x, y in valid_loader:
```

```
    image0 = transform(x[0])
```

```
    image1 = transform(x[1])
```

```
    image2 = transform(x[2])
```

```
    image3 = transform(x[3])
```

```
    image4 = transform(x[4])
```

```
    image5 = transform(x[5])
```

```
    print(x[0].shape)
```

```
    x = x.to(device, dtype=torch.float32)
```

```
    y = y.to(device, dtype=torch.float32)
```

```
    y_pred = m(x)
```

```
    img = y_pred.cpu().detach().numpy()
```

```
    plt.figure(figsize=(30,8))
```

```
    print(img.shape)
```

```
    #subplot(r,c) provide the no. of rows and columns
```

```
    f, axarr = plt.subplots(6,3)
```

```
    f.set_size_inches(12, 30, forward=True)
```

```
    axarr[0,0].imshow(image0)
```

```
    axarr[0,1].imshow(np.squeeze(y.cpu().detach().numpy())[0], cmap='gray')
```

```
    axarr[0,2].imshow(np.squeeze(img)[0], cmap='gray')
```

```
    axarr[1,0].imshow(image1)
```

```
    axarr[1,1].imshow(np.squeeze(y.cpu().detach().numpy())[1], cmap='gray')
```



```
axarr[1,2].imshow(np.squeeze(img)[1], cmap='gray')


axarr[5,0].imshow(image2)
axarr[5,1].imshow(np.squeeze(y.cpu().detach().numpy())[2], cmap='gray')
axarr[5,2].imshow(np.squeeze(img)[2], cmap='gray')

axarr[2,0].imshow(image3)
axarr[2,1].imshow(np.squeeze(y.cpu().detach().numpy())[3], cmap='gray')
axarr[2,2].imshow(np.squeeze(img)[3], cmap='gray')

axarr[3,0].imshow(image4)
axarr[3,1].imshow(np.squeeze(y.cpu().detach().numpy())[4], cmap='gray')
axarr[3,2].imshow(np.squeeze(img)[4], cmap='gray')

axarr[4,0].imshow(image5)
axarr[4,1].imshow(np.squeeze(y.cpu().detach().numpy())[5], cmap='gray')
axarr[4,2].imshow(np.squeeze(img)[5], cmap='gray')
break
```

```
torch.Size([3, 256, 256])
(8, 1, 256, 256)
<Figure size 2160x576 with 0 Axes>
```



```
dice_BCE = evaluate(m, valid_loader, loss_fn, device)
```



```
dice_BCE
```

```
0.4773574968179067
```



```
def dice_score(model, loader, loss_fn, device):
    epoch_loss = 0.0
```

```
    model.eval()
```

```
    with torch.no_grad():
```

```
        for x, y in loader:
```

```
            x = x.to(device, dtype=torch.float32)
```

```
            y = y.to(device, dtype=torch.float32)
```

```
#            x,y=x.type(torch.FloatTensor),y.type(torch.FloatTensor)
```

```
            y_pred = model(x)
```

```
            inputs = y_pred.view(-1)
```

```
            targets = y.view(-1)
```

```
            smooth = 1
```

```
            intersection = (inputs * targets).sum()
```

```
            dice_score = (2.*intersection + smooth)/(inputs.sum() + targets.sum() + smooth)
```

```
            epoch_loss += dice_score
```

```
    epoch_loss = epoch_loss/len(loader)
```

```
    return epoch_loss
```

```
diceScore = dice_score(m, valid_loader, loss_fn, device)
```



```
diceScore.item()
```

```
0.8098794221878052
```



```

class LoadPredictionData(Dataset):
    def __init__(self, images_path):
        super().__init__()

        self.images_path = images_path
        # self.masks_path = masks_path
        self.len = len(images_path)
        self.transform = transforms.Resize(size)

    def __getitem__(self, idx):
        img = Image.open(self.images_path[idx])
        img = self.transform(img)
        img = np.transpose(img, (2, 0, 1))
        img = img/255.0
        img = torch.tensor(img)

        # mask = Image.open(self.masks_path[idx]).convert('L')
        # mask = self.transform(mask)
        # mask = np.expand_dims(mask, axis=0)
        # mask = mask/255.0
        # mask = torch.tensor(mask)

        return img

    def __len__(self):
        return self.len

pred_img = sorted(glob.glob('/content/drive/MyDrive/satellite/download.jpg'))

pred_dataset = LoadPredictionData(pred_img)

len(pred_dataset)

0

prediction_loader = DataLoader(
    dataset=pred_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=2,
)

transform = transforms.ToPILImage()
pred = []
for x in prediction_loader:
    image0 = transform(x[0])

    print(x[0].shape)
    x = x.to(device, dtype=torch.float32)
    # y = y.to(device, dtype=torch.float32)

```

```

y_pred = m(x)
img = y_pred.cpu().detach().numpy()
plt.figure(figsize=(30,8))
print(img.shape)
#subplot(r,c) provide the no. of rows and columns
f, axarr = plt.subplots(6,3)
f.set_size_inches(12, 30, forward=True)
axarr[0,0].imshow(image0)
# axarr[0,1].imshow(np.squeeze(y.cpu().detach().numpy())[0], cmap='gray')
axarr[0,2].imshow(img[0,0,:,:], cmap='gray')

# axarr[1,0].imshow(image1)
# axarr[1,1].imshow(np.squeeze(y.cpu().detach().numpy())[1], cmap='gray')
# axarr[1,2].imshow(np.squeeze(img)[1], cmap='gray')

# axarr[5,0].imshow(image2)
# axarr[5,1].imshow(np.squeeze(y.cpu().detach().numpy())[2], cmap='gray')
# axarr[5,2].imshow(np.squeeze(img)[2], cmap='gray')

# axarr[2,0].imshow(image3)
# axarr[2,1].imshow(np.squeeze(y.cpu().detach().numpy())[3], cmap='gray')
# axarr[2,2].imshow(np.squeeze(img)[3], cmap='gray')

# axarr[3,0].imshow(image4)
# axarr[3,1].imshow(np.squeeze(y.cpu().detach().numpy())[4], cmap='gray')
# axarr[3,2].imshow(np.squeeze(img)[4], cmap='gray')

# axarr[4,0].imshow(image5)
# axarr[4,1].imshow(np.squeeze(y.cpu().detach().numpy())[5], cmap='gray')
# axarr[4,2].imshow(np.squeeze(img)[5], cmap='gray')
break

```

!pip install gradio

```

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels
Collecting gradio
  Downloading gradio-3.0.10-py3-none-any.whl (5.1 MB)
    |████████████████████████████████████████| 5.1 MB 8.0 MB/s
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages
Collecting ffmpeg
  Downloading ffmpeg-0.3.0.tar.gz (4.8 kB)
Collecting pycryptodome
  Downloading pycryptodome-3.14.1-cp35-abi3-manylinux2010_x86_64.whl (2.0 MB)
    |████████████████████████████████████████| 2.0 MB 53.0 MB/s
Collecting uvicorn
  Downloading uvicorn-0.17.6-py3-none-any.whl (53 kB)
    |████████████████████████████████████████| 53 kB 2.6 MB/s
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages
Collecting analytics-python
  Downloading analytics_python-1.4.0-py2.py3-none-any.whl (15 kB)
Collecting python-multipart
  Downloading python-multipart-0.0.5.tar.gz (32 kB)
Collecting aiohttp
  Downloading aiohttp-3.8.1-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.many
    |████████████████████████████████████████| 1.1 MB 44.1 MB/s
Collecting orjson

```

```

Downloading orjson-3.7.0-cp37-cp37m-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
| 256 kB 68.0 MB/s
Requirement already satisfied: pandas in /usr/local/lib/python3.7/dist-packages (1.3.4)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.7/dist-packages (3.0.1)
Collecting paramiko
  Downloading paramiko-2.11.0-py2.py3-none-any.whl (212 kB)
  | 212 kB 67.8 MB/s
Collecting markdown-it-py[linkify,plugins]
  Downloading markdown_it_py-2.1.0-py3-none-any.whl (84 kB)
  | 84 kB 3.9 MB/s
Requirement already satisfied: pillow in /usr/local/lib/python3.7/dist-packages (9.0.1)
Collecting pydub
  Downloading pydub-0.25.1-py2.py3-none-any.whl (32 kB)
Collecting fastapi
  Downloading fastapi-0.78.0-py3-none-any.whl (54 kB)
  | 54 kB 3.5 MB/s
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (1.21.0)
Collecting aiosignal>=1.1.2
  Downloading aiosignal-1.2.0-py3-none-any.whl (8.2 kB)
Requirement already satisfied: charset-normalizer<3.0,>=2.0 in /usr/local/lib/python3.7/dist-packages (2.0.12)
Collecting yarl<2.0,>=1.0
  Downloading yarl-1.7.2-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2014_x86_64.whl
  | 271 kB 56.5 MB/s
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.7/dist-packages (21.2.0)
Collecting multidict<7.0,>=4.5
  Downloading multidict-6.0.2-cp37-cp37m-manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux2014_x86_64.whl
  | 94 kB 4.4 MB/s
Collecting frozenlist>=1.1.1
  Downloading frozenlist-1.3.0-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2014_x86_64.whl
  | 144 kB 59.2 MB/s
Requirement already satisfied: typing-extensions>=3.7.4 in /usr/local/lib/python3.7/dist-packages (4.1.1)
Collecting async-timeout<5.0,>=4.0.0a3
  Downloading async_timeout-4.0.2-py3-none-any.whl (5.8 kB)
Collecting asyncctest==0.13.0
  Downloading asyncctest-0.13.0-py3-none-any.whl (26 kB)

```

```

def segment(image):
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    cv2.imwrite('/content/drive/MyDrive/satellite/test.jpg',image)
    pred_img = sorted(glob.glob('/content/drive/MyDrive/satellite/test.jpg'))
    print(image.shape)
    pred_dt = LoadPredictionData(pred_img)
    print(len(pred_dt))
    prediction_lr = DataLoader(dataset=pred_dt, batch_size=batch_size, shuffle=False, num_workers=1)
    print(len(prediction_lr))
    for x in prediction_lr:
        print("#####")
        print(x[0].shape)
        image0 = transform(x[0])
        print("#####")
        x = x.to(device, dtype=torch.float32)
        print("#####")
        y_pred = m(x)
        print("#####")
        img = y_pred.cpu().detach().numpy()
        break
    print("#####")

```

```

return img[0,0,:,:]
# return

```

```

# ans=segment(pred_img)
# print(ans.shape)
# plt.figure(figsize=(9,9))
# plt.subplot(1,1,1)
# plt.imshow(ans,cmap='gray')

```

```
import numpy as np
```

```

import gradio as gr
image = gr.inputs.Image(shape = (256, 256),type="numpy")
demo = gr.Interface(fn=segment, inputs=image, outputs="image")
demo.launch(auth=('user', 'password'),debug=True)

```

```

/usr/local/lib/python3.7/dist-packages/gradio/deprecation.py:40: UserWarning: `option` is deprecated. Use `kwargs` instead.
warnings.warn(value)

```

Colab notebook detected. This cell will run indefinitely so that you can see errors & warnings. [Running on public URL: https://50712.gradio.app](https://50712.gradio.app)

This share link expires in 72 hours. For free permanent hosting, check out Spaces (<https://github.com/openai/gpt-4>)

```

1
1
#####
torch.Size([3, 256, 256])
#####
#####
#####
*****
(256, 256, 3)
1
1
#####
torch.Size([3, 256, 256])
#####
#####
#####
*****
(256, 256, 3)
1
1
#####
torch.Size([3, 256, 256])
#####
#####
#####
*****
(256, 256, 3)
1
1
#####
torch.Size([3, 256, 256])
#####
#####
#####

```

\*\*\*\*\*

Keyboard interruption in main thread... closing server.

(<gradio.routes.App at 0x7f2ee3ecd410>,

'<http://127.0.0.1:7860/>',

'<https://50712.gradio.app>'')

